

# DTCC

## 数 / 造 / 未 / 来

### 第十二届中国数据库技术大会

DATABASE TECHNOLOGY CONFERENCE CHINA 2021



2021 年 10 月 18 日 - 20 日 | 北京国际会议中心





# 摸着Oracle过河 ----

## 大幅提升PostgreSQL性能分享

吕海波 ( VAGE )




# 主题：

## 优化PostgreSQL Partial Writes（页裂）体系

### 大幅提升整体性能



# 页分裂：部分写，Partial Writes



## ➤ 八仙过海 各显神通

- MySQL解决方案：Double Write
- Oracle解决方案：众说纷纭（高端硬件说、自有原子写说）
- PG解决方案：Full Page Writes（简称FPW）





# 页分裂：部分写，Partial Writes



DTCC 2021  
第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021

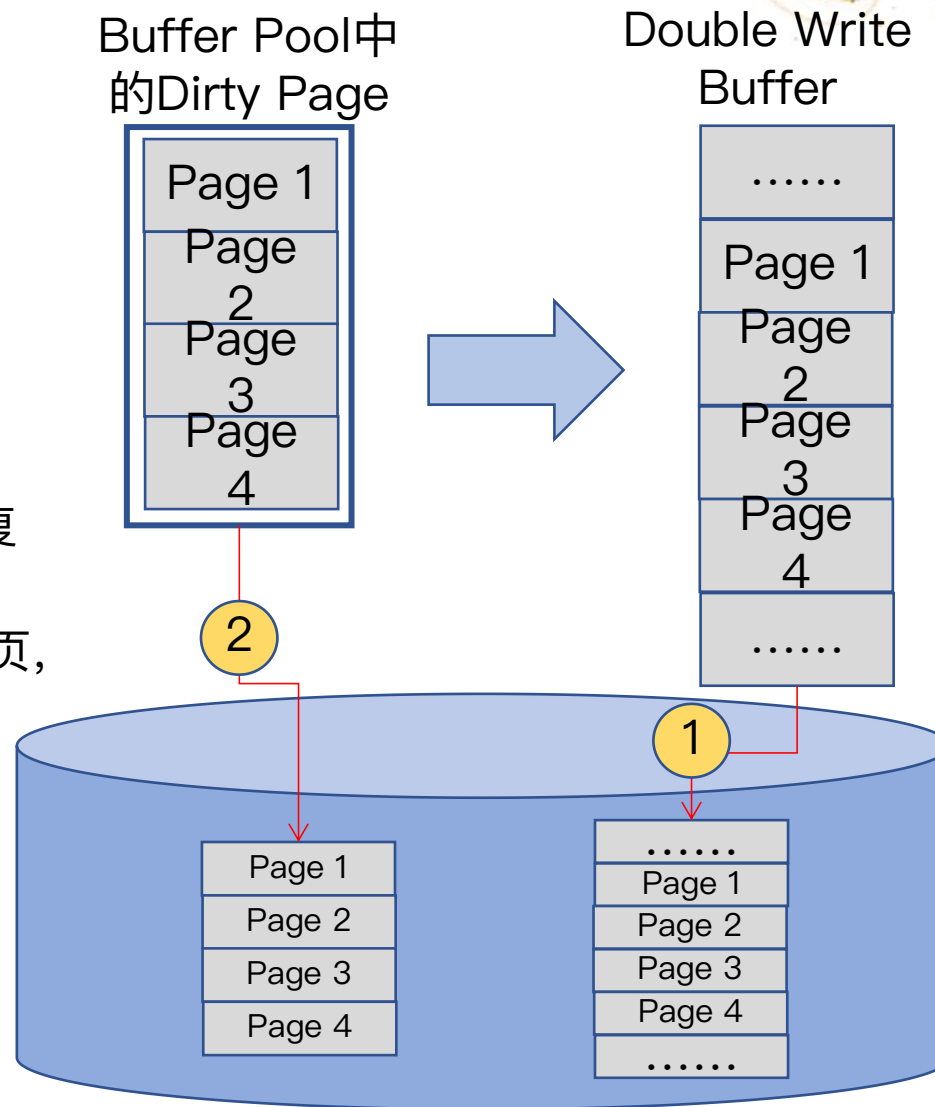
## ➤ MySQL Double Write

先写Double Write Buffer

再写Buffer Pool中的脏页

脏页写失败了，使用Double Write中的页进行恢复

Double Write写失败了，不写Buffer Pool中的脏页，脏页直接丢弃，依赖BinLog恢复。



数，造，未，来



IT168.com

ChinaUnix.net

ITPUB

# 页分裂：部分写，Partial Writes



DTCC 2021  
第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021

## ➤ MySQL Double Write

先写Double Write Buffer

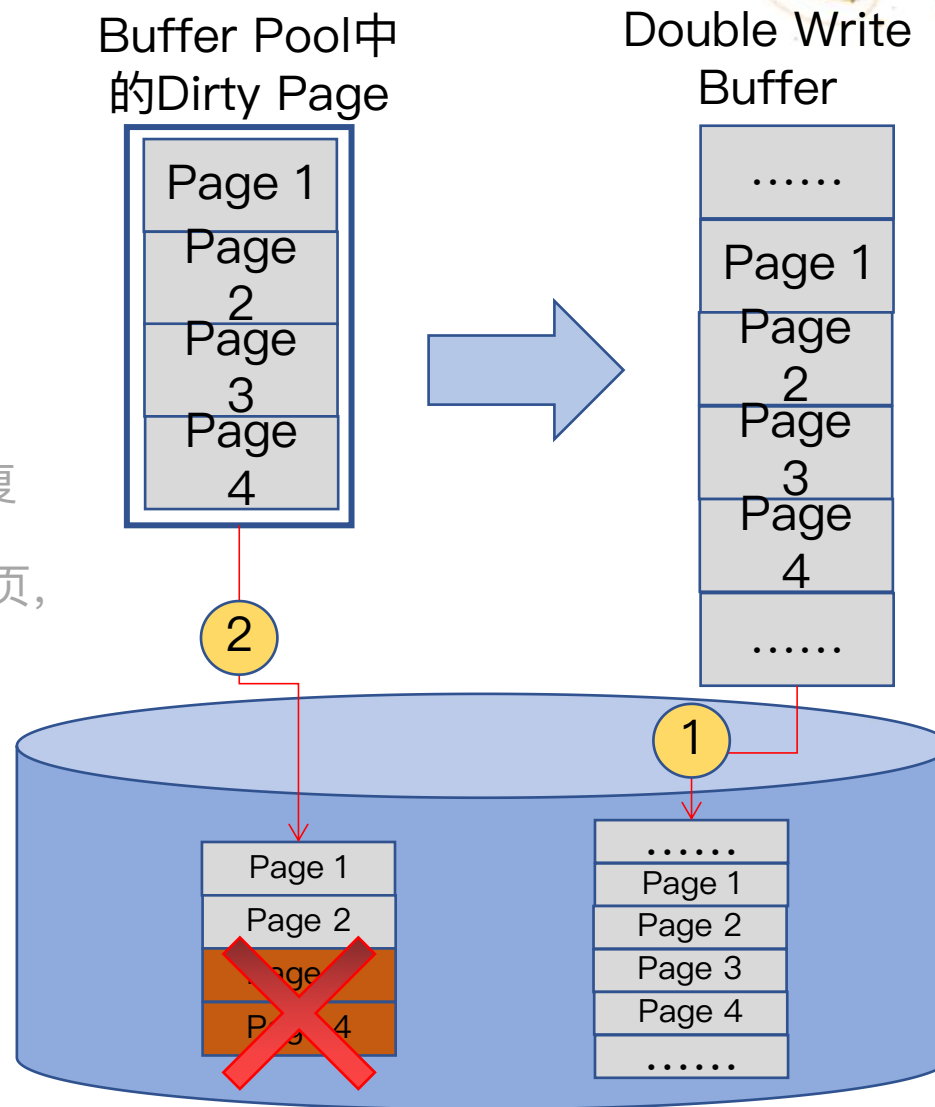
再写Buffer Pool中的脏页

脏页写失败了，使用Double Write中的页进行恢复

Double Write写失败了，不写Buffer Pool中的脏页，脏页直接丢弃，依赖BinLog恢复。

为什么MySQL必须要依赖双写？

为什么不能使用BinLog进行恢复？



数，造，未，来



# 页分裂：部分写，Partial Writes



DTCC 2021  
第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021

## ➤ MySQL Double Write

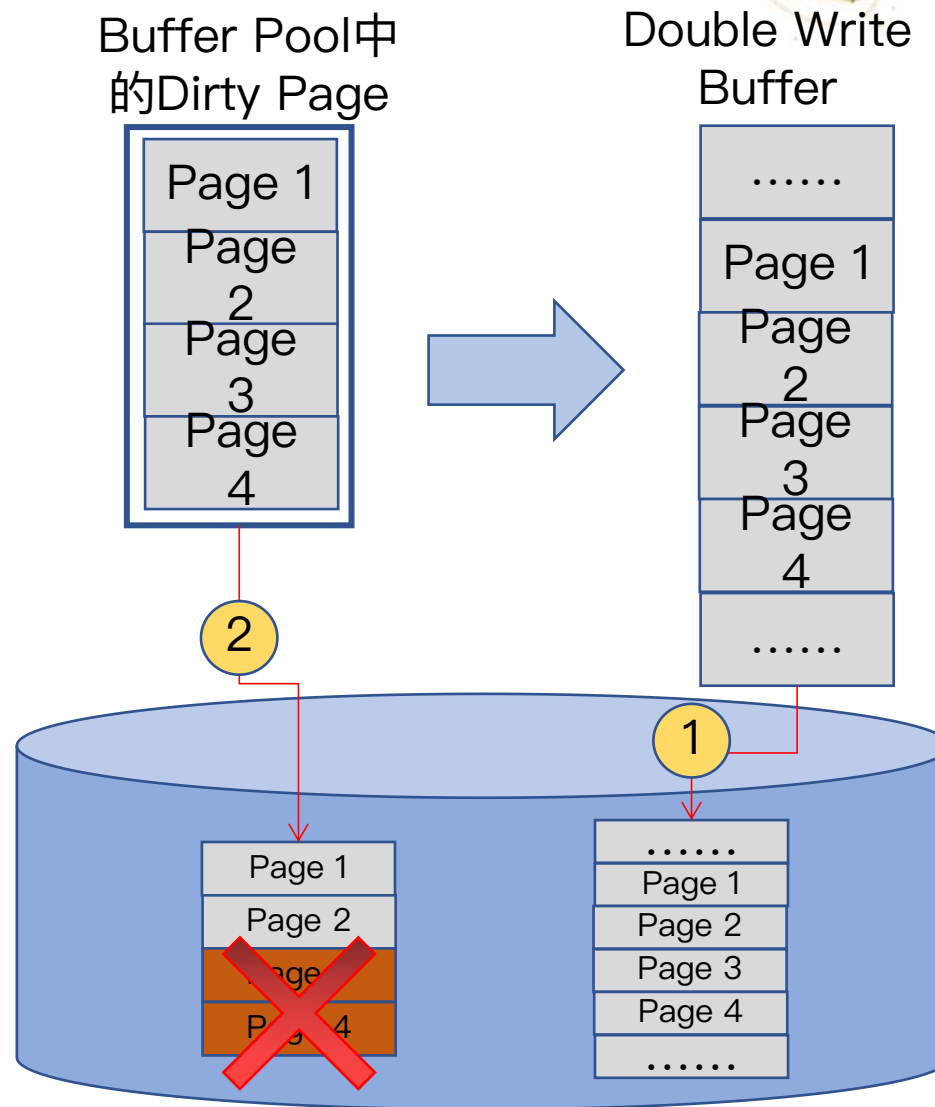
以Insert为例：

脏页中是Insert的新行，

BinLog记录的是SQL（逻辑日志），

使用BinLog恢复相当于再次执行Insert，

因此，



数，造，未，来



IT168.com

ChinaUnix.net

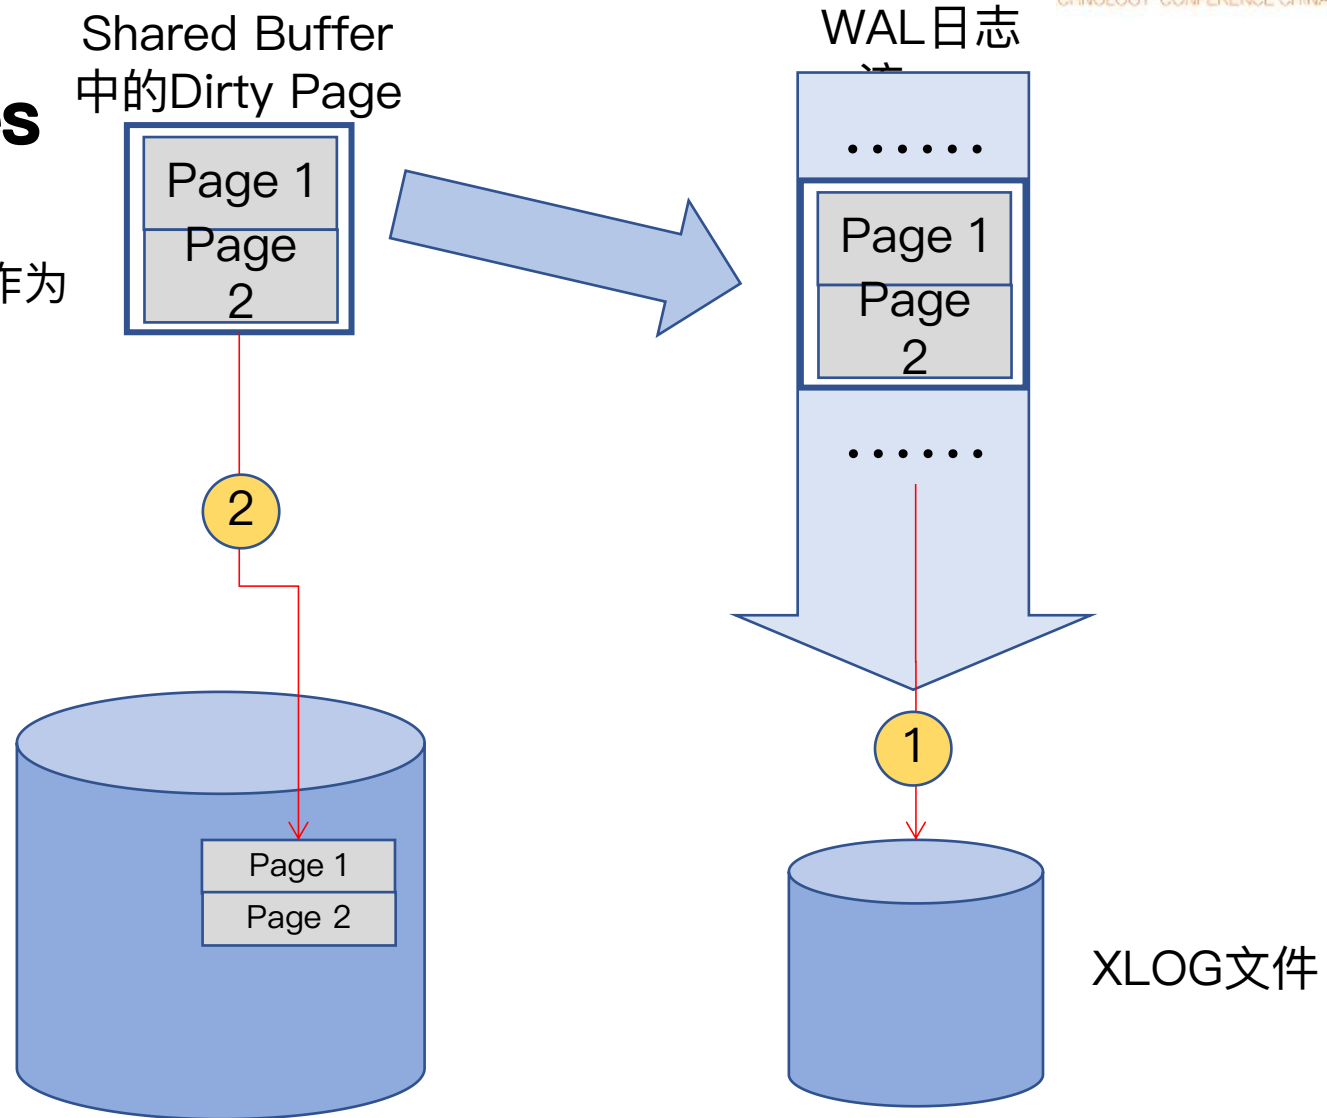
ITPUB



# 页分裂：部分写，Partial Writes

## ➤ PostgreSQL Full Page Writes

每次Checkpoint后，首次修改页，拷贝整页到日志中作为备份。

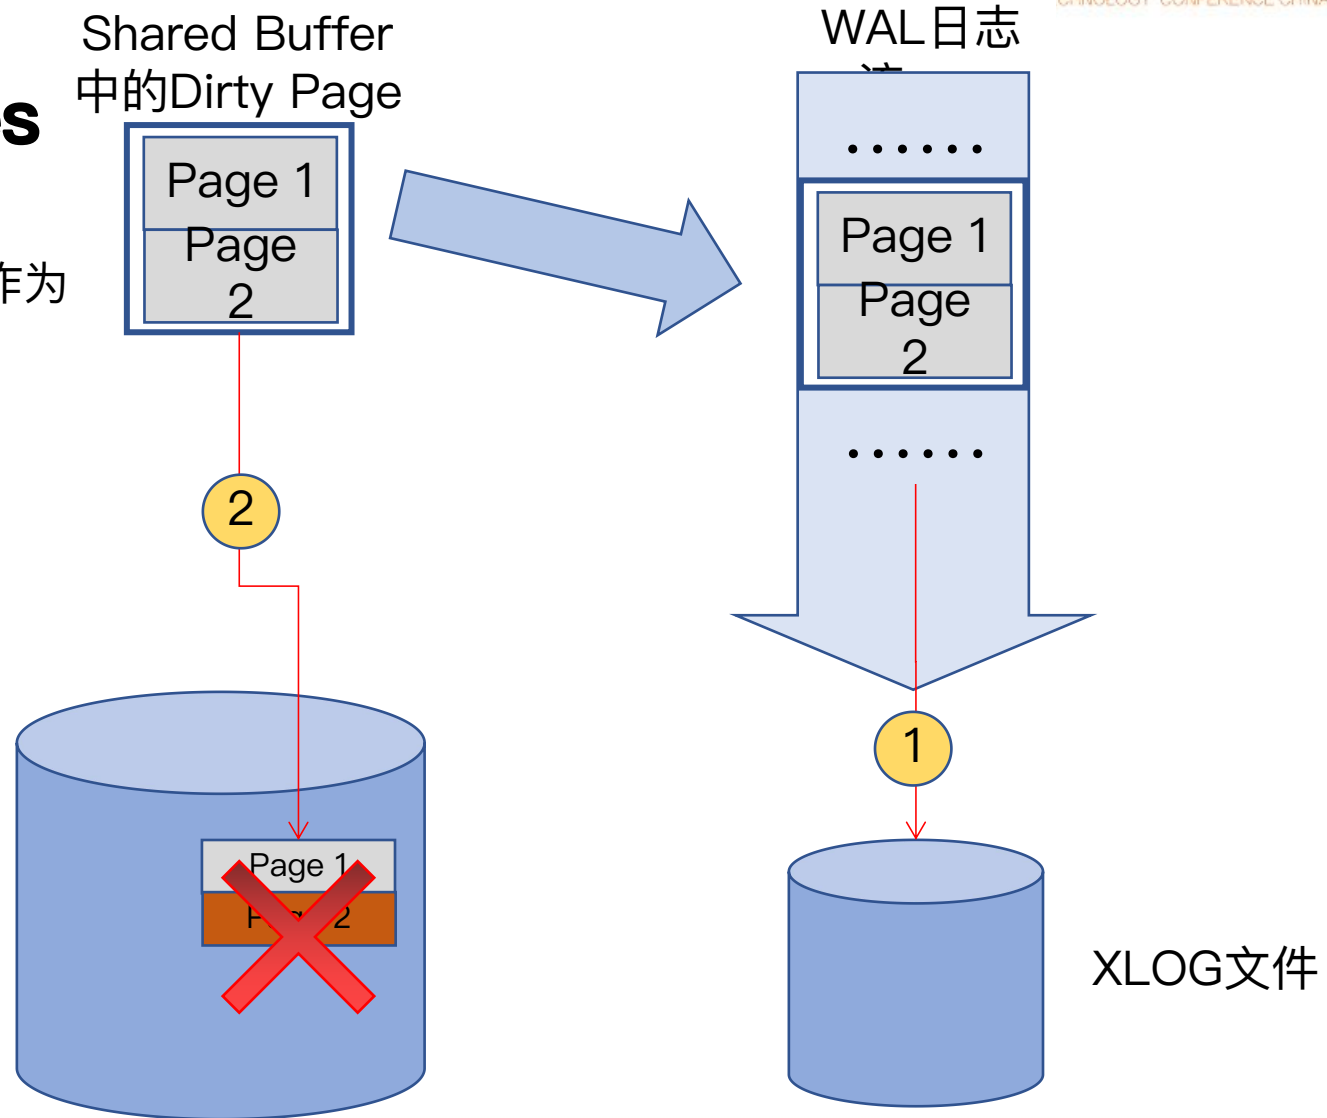


# 页分裂：部分写，Partial Writes

## ➤ PostgreSQL Full Page Writes

每次Checkpoint后，首次修改页，拷贝整页到日志中作为备份。

出现Partial Writes，以此备份为基础，对页进行恢复。



# 页分裂：部分写，Partial Writes

## ➤ 性能影响对比，MySQL Double Write性能影响：

```
OLTP test statistics:
  queries performed:
    read:          3165806
    write:         904516
    other:         452258
    total:         4522580
  transactions:    226129 (1254.75 per sec.)
  deadlocks:      0 (0.00 per sec.)
  read/write requests: 4070322 (22585.49 per sec.)
  other operations: 452258 (2509.50 per sec.)

General statistics:
  total time:      180.2185s
  total number of events: 226129
  total time taken by event execution: 5762.8595s
  response time:
    min:          6.98ms
    avg:          25.48ms
    max:          631.13ms
    approx. 95 percentile: 38.26ms

Threads fairness:
  events (avg/stddev): 7066.5312/25.07
  execution time (avg/stddev): 180.0894/0.07

[mysql@localhost ~]$
```

```
OLTP test statistics:
  queries performed:
    read:          3249344
    write:         928384
    other:         464192
    total:         4641920
  transactions:    232096 (1289.34 per sec.)
  deadlocks:      0 (0.00 per sec.)
  read/write requests: 4177728 (23208.05 per sec.)
  other operations: 464192 (2578.67 per sec.)

General statistics:
  total time:      180.0120s
  total number of events: 232096
  total time taken by event execution: 5759.5750s
  response time:
    min:          7.30ms
    avg:          24.82ms
    max:          561.23ms
    approx. 95 percentile: 37.46ms

Threads fairness:
  events (avg/stddev): 7253.0000/31.00
  execution time (avg/stddev): 179.9867/0.00

[mysql@localhost ~]$
```





# 页分裂：部分写，Partial Writes



**DTCC 2021**  
第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021

## ➤ 性能影响对比，PostgreSQL Full Page Writes性能影响：

```
OLTP test statistics:
queries performed:
  read:                2627954
  write:               750844
  other:               375422
  total:               3754220
transactions:         187711 (1041.86 per sec.)
deadlocks:            0      (0.00 per sec.)
read/write requests:  3378798 (18753.42 per sec.)
other operations:     375422 (2083.71 per sec.)

General statistics:
total time:           180.1697s
total number of events: 187711
total time taken by event execution: 11522.4063s
response time:
  min:                 2.69ms
  avg:                 61.38ms
  max:                 1206.33ms
  approx. 95 percentile: 197.71ms

Threads fairness:
events (avg/stddev):  2932.9844/30.51
execution time (avg/stddev): 180.0376/0.04

[postgres@localhost data]$
```

```
OLTP test statistics:
queries performed:
  read:                5875492
  write:               1678712
  other:               839356
  total:               8393560
transactions:         419678 (2331.48 per sec.)
deadlocks:            0      (0.00 per sec.)
read/write requests:  7554204 (41966.55 per sec.)
other operations:     839356 (4662.95 per sec.)

General statistics:
total time:           180.0053s
total number of events: 419678
total time taken by event execution: 11516.2816s
response time:
  min:                 2.50ms
  avg:                 27.44ms
  max:                 650.34ms
  approx. 95 percentile: 51.67ms

Threads fairness:
events (avg/stddev):  6557.4688/43.48
execution time (avg/stddev): 179.9419/0.01
```



数，造，未，来

Hadoop

OLAP

HANA

MPP

DB2

Oracle

SQL

IT168.com

ChinaUnix.net

ITPUB



走两步





- 拦截系统调用，修改传入OS内核的参数



- 拦截系统调用，修改传入OS内核的参数
- Linux, systemtap



# Partial Writes模拟：MySQL篇

➤ I/O相关系统函数取决于innodb\_use\_native\_aio参数

on: Sys\_io\_submit

off: pwrite



# Partial Writes模拟：MySQL篇



DTCC 2021  
第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021

在MySQL中，创建测试表，插入200行数据：

```
CREATE TABLE `vage` (  
  `id1` int(11) NOT NULL,  
  `id2` int(11) DEFAULT NULL,  
  `c1` varchar(100) DEFAULT NULL,  
  `c2` varchar(100) DEFAULT NULL,  
  PRIMARY KEY (`id1`)  
) ENGINE=InnoDB DEFAULT CHARSET=utf8  
  
delimiter $$  
CREATE PROCEDURE vage_init(in nums int)  
BEGIN  
  declare done int default 0;  
  declare i int;  
  declare v_id1 int;  
  
  set i=0;  
  while i<=nums DO  
    insert into vage values(i, i+100, concat('AAAAAA', i), concat('BBBBBBBBB', i));  
    set i=i+1;  
  end while;  
  commit;  
END$$  
delimiter ;  
  
call vage_init(200);
```



数，造，未，来





# Partial Writes模拟：MySQL篇

MySQL是从上到下使用页中空间，因此数据行要多一些，要保证数据超过4K：

```
0000bfff 00 00 00 00 00 00 00 00 0f 65 b1 0c 00 29 f7 19 | .....e...)..
0000c000 5c 3b 29 d9 00 00 00 03 ff ff ff ff ff ff ff | \;).....
0000c010 00 00 00 00 00 2b b4 92 45 bf 00 00 00 00 00 00 | .....+..E.....
0000c020 00 00 00 00 00 17 00 33 26 15 80 cb 00 00 00 00 | .....3&.....
0000c030 25 ec 00 02 00 c8 00 c9 00 00 00 00 00 00 00 00 | %.....
0000c040 00 00 00 00 00 00 00 00 00 29 00 00 00 17 00 00 | .....).....
0000c050 00 02 00 f2 00 00 00 17 00 00 00 02 00 32 01 00 | .....2.....
0000c060 02 00 1d 69 6e 66 69 6d 75 6d 00 06 00 0b 00 00 | ...infimum.....
0000c070 73 75 70 72 65 6d 75 6d 09 07 00 00 00 10 00 2d | supremum.....-
0000c080 80 00 00 00 00 00 00 00 05 07 a7 00 00 01 1b 01 | .....
0000c090 10 80 00 00 64 41 41 41 41 41 41 30 42 42 42 42 | ....dAAAAA0BBBB
0000c0a0 42 42 42 42 30 09 07 00 00 00 18 00 2d 80 00 00 | BBBB0.....-...
0000c0b0 01 00 00 00 00 05 08 a8 00 00 01 1c 01 10 80 00 | .....
0000c0c0 00 65 41 41 41 41 41 41 31 42 42 42 42 42 42 | .eAAAAAA1BBBBBBB
0000c0d0 42 31 09 07 00 00 00 20 00 2d 80 00 00 02 00 00 | B1.....-.....
0000c0e0 00 00 05 0b aa 00 00 01 1e 01 10 80 00 00 66 41 | .....fA
0000c0f0 41 41 41 41 41 32 42 42 42 42 42 42 42 32 09 | AAAAA2BBBBBBBB2.
0000c100 07 00 04 00 28 00 2d 80 00 00 03 00 00 00 05 | ....(-.....
```

使用“dd if=vage.ibd bs=16384 | hexdump -C |more”观察表对应文件，从0xc000开始，是表数据对应的页（之前数据是表头）。

两百行数据，目前占据一个MySQL页，即从0xC000到0x10000。





# Partial Writes模拟：MySQL篇

MySQL是从上到下使用页中空间，因此数据行要多一些，要保证数据超过4K：

0000e5a0	41	41	41	41	41	31	39	38	42	42	42	42	42	42	42	AAAAA198BBBBBBBB
0000e5b0	31	39	38	0b	09	00	00	06	48	00	31	80	00	00	c7	198.....H.1.....
0000e5c0	00	00	00	05	d2	d0	00	00	01	4d	01	10	80	00	01	.....M.....+
0000e5d0	41	41	41	41	41	41	31	39	39	42	42	42	42	42	42	AAAAAA199BBBBBBBB
0000e5e0	42	31	39	39	0b	09	00	00	06	50	da	84	80	00	00	B199.....P.....
0000e5f0	00	00	00	00	05	d3	d1	00	00	01	4e	01	10	80	00	.....N.....
0000e600	2c	41	41	41	41	41	32	30	30	42	42	42	42	42	42	AAAAAA200BBBBBB
0000e610	42	42	32	30	30	00	00	00	00	00	00	00	00	00	00	BB200.....
0000e620	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	.....
*																
0000ff90	00	00	00	70	24	f7	24	33	23	6f	22	ab	21	e7	21	...p\$. \$3#o".! !#
0000ffa0	20	5f	1f	9b	1e	d7	1e	13	1d	4f	1c	8b	1b	c7	1b	_. ....O. ....
0000ffb0	1a	3f	19	7b	18	b7	17	f3	17	2f	16	6b	15	a7	14	.?.{...../.k....
0000ffc0	14	1f	13	5b	12	99	11	dd	11	21	10	65	0f	a9	0e	...[.....!.e....
0000ffd0	0e	31	0d	75	0c	b9	0b	fd	0b	41	0a	85	09	c9	09	.1.u.....A.....
0000ffe0	08	51	07	95	06	d9	06	1d	05	61	04	a5	03	e9	03	.Q.....a.....-
0000fff0	02	71	01	bb	01	07	00	63	5c	3b	29	d9	00	2b	b4	.q.....c\;)..+..
00010000	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	.....
*																

最后一行，开始自0xe5e0这一行。我们以最后一行的c1列为目标，它在0xe601处。

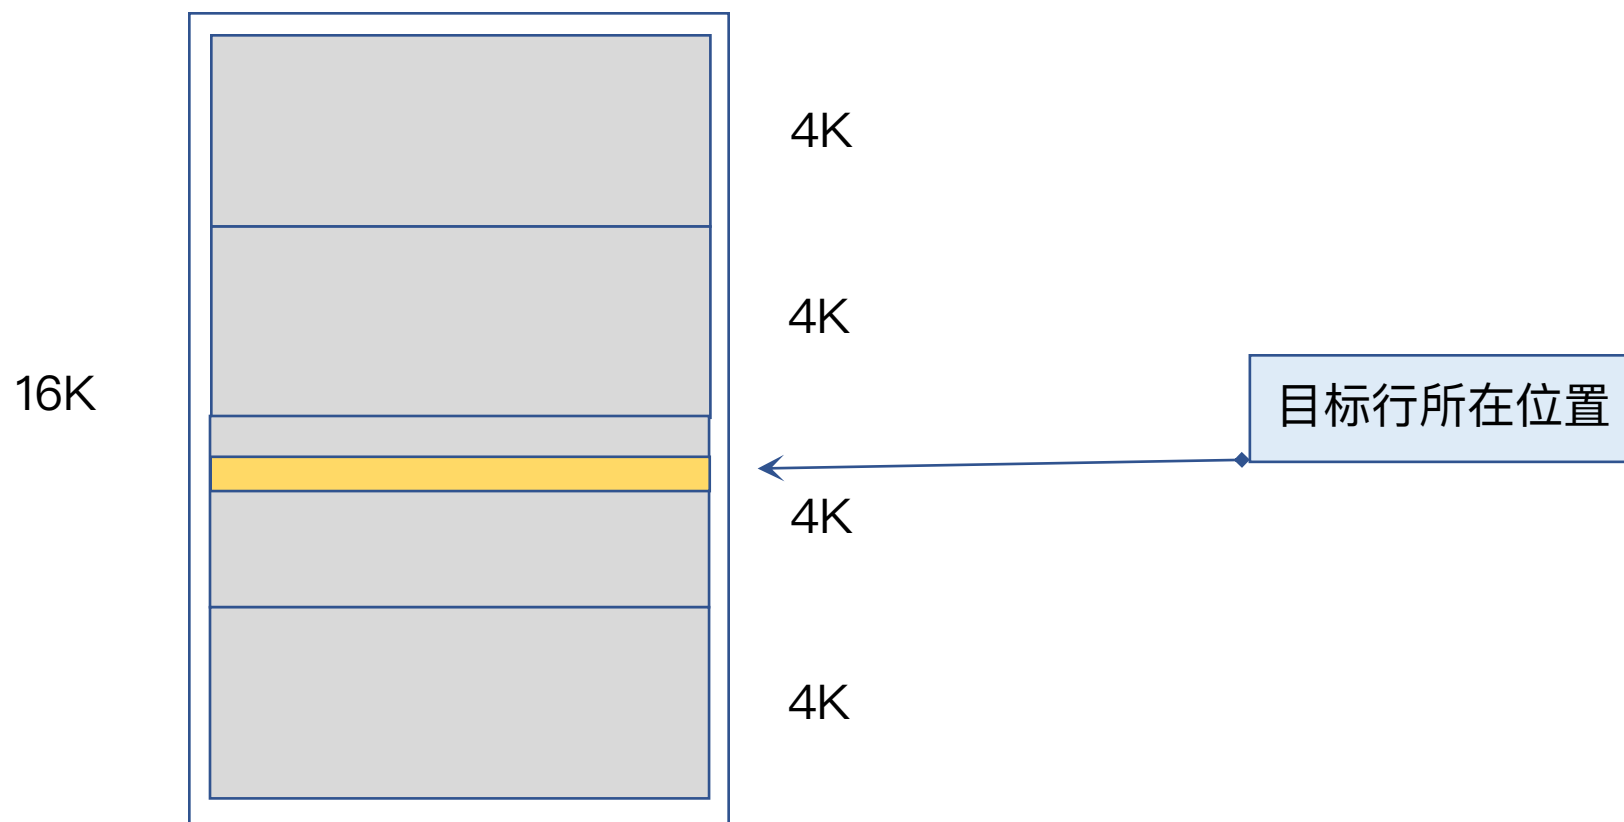


# Partial Writes模拟：MySQL篇



DTCC 2021  
第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021

MySQL是从上到下使用页中空间，因此数据行要多一些，要保证数据超过4K：



数，造，未，来




IT168.com

ChinaUnix.net

ITPUB

# Partial Writes模拟：MySQL篇



```
NAME
    pread, pwrite - read from or write to a file descriptor at a given offset
SYNOPSIS
    #include <unistd.h>

    ssize_t pread(int fd, void *buf, size_t count, off_t offset);
    ssize_t pwrite(int fd, const void *buf, size_t count, off_t offset);
Feature Test Macro Requirements for glibc (see feature_test_macros(7)):


    pread(), pwrite():
        _XOPEN_SOURCE >= 500
        || /* Since glibc 2.12: */ _POSIX_C_SOURCE >= 200809L
```

发出修改目标行的SQL，触发对目标页的写操作，拦截OS I/O系统调用，修改count参数。





# Partial Writes模拟：MySQL篇



```
NAME
    pread, pwrite - read from or write to a file descriptor at a given offset
SYNOPSIS
    #include <unistd.h>
    ssize_t pread(int fd, void *buf, size_t count, off_t offset);
    ssize_t pwrite(int fd, const void *buf, size_t count, off_t offset);
    Feature Test Macro Requirements for glibc (see feature_test_macros(7)):
    pread(), pwrite():
        _XOPEN_SOURCE >= 500
        || /* Since glibc 2.12: */ _POSIX_C_SOURCE >= 200809L
```

count参数原来是16384，将其改为8192。造成一个页，只写了一半的效果，也就是页的Partial Writes。同时还要返回I/O错误，让MySQL知道，I/O并没有成功完成。因为突然的OS宕机，也是同样的情况：

- 只完成部分I/O
- 最终I/O错误



# Partial Writes模拟：MySQL篇



DTCC 2021  
第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021

```
.....  
probe syscall.pwrite  
{
```

```
if (execname() = mysql>
```

```
mysql> update vage set c1=lower(c1) where id1=200;  
Query OK, 1 row affected (0.01 sec)  
Rows matched: 1  Changed: 1  Warnings: 0
```

```
{  
    printf("%04d %s %s %x %x %x %x\n", tid(), pp(), pfunc(), $fd, $buf, $count, $pos)  
}  
}  
.....
```

```
[root@localhost iotest]# stap -g dbpage_m.stp mysql
```

```
Begin...
```

```
14658 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 4 7fd0a4495200 600 2ba800
```

```
return:14658 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
```

```
14643 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 4 7fd0a5495200 200 2bac00
```

```
return:14643 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
```

```
14638 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 b 7fd0a052c000 18000 100000
```

```
return:14638 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
```

```
14635 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 b 7fd0a6f20000 4000 508000
```

```
14634 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 b 7fd0a6ce4000 4000 0
```

```
14637 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 b 7fd0a6b90000 4000 380000
```

```
return:14634 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
```

```
14634 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 28 7fd0a6ef8000 4000 c000
```

```
14635 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 b 7fd0a7018000 4000 5fc000
```

```
return:14637 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
```

```
14634 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 b 7fd0a6a24000 4000 14000
```

```
return:14634 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
```

```
14643 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 4 7fd0a4495200 200 2bac00
```

```
return:14643 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
```

```
14629 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 4 2a69a00 200 200
```

```
return:14629 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
```



数，造，未，来

Hadoop

OLAP

HANA

MPP

Oracle

SQL

IT168.com

ChinaUnix.net

ITPUB



# Partial Writes模拟：MySQL篇



DTCC 2021  
第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021

.....  
printf("%d %s %s %x %x %x %x\n", tid(), pp(), ppfunc(), \$fd, \$buf, \$count, \$pos)  
.....

```
[root@localhost iotest]# stap -g dbpage_m.stp mysqld
Begin...
14658 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 4 7fd0a4495200 600 2ba800
return:14658 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
14643 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 4 7fd0a5495200 200 2bac00
return:14643 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64

14638 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 b 7fd0a052c000 18000 100000
return:14638 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
14635 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 b 7fd0a6f20000 4000 508000
14634 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 b 7fd0a6ce4000 4000 0
14637 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 b 7fd0a6b90000 4000 380000
return:14634 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
14634 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 28 7fd0a6ef8000 4000 c000
14635 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 b 7fd0a7018000 4000 5fc000
return:14637 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
14634 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 b 7fd0a6a24000 4000 14000
return:14634 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
14643 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 4 7fd0a4495200 200 2bac00
return:14643 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
14629 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 4 2a69a00 200 200
return:14629 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
```

目标表的更新在中间，fd是0x28（十进制为40），它是表对应的文件：vage.ibd。  
可以在/proc/(mysqld进程号)/fd目录中确认此点。

count参数为0x4000（16384），pos参数为0xc000。



数，造，未，来





# Partial Writes模拟：MySQL篇



DTCC 2021  
第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021

```
.....  
printf("%d %s %s %x %x %x %x\n", tid(), pp(), ppfunc(), $fd, $buf, $count, $pos)  
.....
```

```
[root@localhost iotest]# stap -g dbpage_m.stp mysql  
Begin...  
14658 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 4 7fd0a4495200 600 2ba800  
return:14658 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64  
14643 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 4 7fd0a5495200 200 2bac00  
return:14643 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64  
  
14638 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 b 7fd0a052c000 18000 100000  
return:14638 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64  
14635 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 b 7fd0a6f20000 4000 508000  
14634 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 b 7fd0a6ce4000 4000 0  
14637 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 b 7fd0a6b90000 4000 380000  
return:14634 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64  
14634 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 28 7fd0a6ef8000 4000 c000  
14635 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 b 7fd0a7018000 4000 5fc000  
return:14637 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64  
14634 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 b 7fd0a6a24000 4000 14000  
return:14634 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64  
14643 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 4 7fd0a4495200 200 2bac00  
return:14643 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64  
14629 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 4 2a69a00 200 200  
return:14629 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
```

对0xb号文件（系统表对应文件：ibdata1）有4次写操作。其中第一次写0x18000（6个页）字节的，是写Double Write Buffer。可以使用dd 命令，读出偏移0x100,000的数据进行确认。

后面几个页，是UNDO和其他内容。



数，造，未，来





# Partial Writes模拟：MySQL篇



DTCC 2021  
第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021

```
.....  
printf("%d %s %s %x %x %x %x\n", tid(), pp(), ppfunc(), $fd, $buf, $count, $pos)  
.....
```

```
[root@localhost iotest]# stap -g dbpage_m.stp mysql  
Begin...  
14658 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 4 7fd0a4495200 600 2ba800  
return:14658 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64  
14643 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 4 7fd0a5495200 200 2bac00  
return:14643 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64  
  
14638 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 b 7fd0a052c000 18000 100000  
return:14638 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64  
14635 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 b 7fd0a6f20000 4000 508000  
14634 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 b 7fd0a6ce4000 4000 0  
14637 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 b 7fd0a6b90000 4000 380000  
return:14634 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64  
14634 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 28 7fd0a6ef8000 4000 c000  
14635 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 b 7fd0a7018000 4000 5fc000  
return:14637 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64  
14634 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 b 7fd0a6a24000 4000 14000  
return:14634 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64  
14643 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 4 7fd0a4495200 200 2bac00  
return:14643 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64  
14629 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 4 2a69a00 200 200  
return:14629 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
```

双写相关I/O，仍然保证成功。我们的目标，拦截系统调用，修改参数，使表/索引页的I/O出现Partial Writes。

方法：修改count参数，造成部分写入成功。修改fd参数，使写16K的I/O最终报错。



数，造，未，来



# Partial Writes模拟：MySQL篇



DTCC 2021

第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021

## ➤ 最终的脚本：

修改\$count为0x1000，成功写为4K。

后面的脚本，是为了造成I/O错误。

前面的条件，是跳过双写相关的I/O，保证双写可以成功完成。

```
probe syscall.pwrite
{
    if (execname() == @1)
    {
        printf("old:%d %s %s %x %x %x %x\n", tid(), pp(), ppfunc(), $fd, $buf, $count, $pos)
        if ( ($fd == 0x28 || $fd == 0xb) && ($count == 0x4000 || (($buf - bufaddr) == 0x1000)) )
        {
            $count = 0x1000
            if ( ($buf - bufaddr) == 0x1000 )
            {
                i += 1
                $buf = bufaddr
                $pos = $pos - (i * 0x1000)
                $fd = 0x1000 + $fd
                bufaddr = 0
            }
            bufaddr = $buf
            printf("update value...\n")
        }
        printf("%d %s %s %x %x %x %x\n", tid(), pp(), ppfunc(), $fd, $buf, $count, $pos)
        if ( $fd == 0x28 )
            exit()
    }
}
```



数，造，未，来



IT168.com

ChinaUnix.net

ITPUB



# Partial Writes模拟：MySQL篇



DTCC 2021

第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021

```
[root@localhost iotest]# stap -g dbpage_m.stp mysqld
Begin...
old:14658 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 4 7fd0a4495200 800 2bac00
14658 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 4 7fd0a4495200 800 2bac00
old:14643 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 4 7fd0a5495200 200 2bb200
14643 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 4 7fd0a5495200 200 2bb200
old:14638 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 b 7fd0a052c000 18000 100000
14638 mysql>
old:14638 mysql> update vage set c1=upper(c1) where id1=200;
Query OK, 1 row affected (0.00 sec)
update Rows matched: 1 Changed: 1 Warnings: 0
14636
old:14636 mysql> select * from vage where id1=200;
update ERROR 2006 (HY000): MySQL server has gone away
14637 No connection. Trying to reconnect...
old:14637 ERROR 2002 (HY000): Can't connect to local MySQL server through socket '/tmp/mysql.sock' (2)
update ERROR:
14635 Can't connect to the server
old:14635
update
14634 mysql>
old:14635 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 b 7fd0a6f21000 3000 509000
14635 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 b 7fd0a6f21000 3000 509000
old:14634 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 b 7fd0a6ce5000 3000 1000
update value...
14634 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 100b 7fd0a6ce4000 1000 0
old:14637 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 b 7fd0a6ce5000 3000 205000
```



数，造，未，来

Hadoop

OLAP

HANA

MPP

DB2

Oracle

SQL

IT168.com

ChinaUnix.net

ITPUB



# Partial Writes模拟：MySQL篇

脚本还存在一些问题，并不能使所有表/索引页的修改都是Partial Writes。



# Partial Writes模拟：MySQL篇



DTCC 2021  
第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021

启动MySQL，可以看到页损坏报错，启动失败：

```
2021-10-09T15:20:31.779551+08:00 0 [Note] InnoDB: If the mysqld execution user is authorized, page cleaner thread priority can be changed. See the man page of setpriority().  
2021-10-09T15:20:31.797777+08:00 0 [ERROR] InnoDB: Database page corruption on disk or a failed file read of page [page id: space=0, page number=5]. You may have to recover from a backup.  
2021-10-09T15:20:31.797831+08:00 0 [Note] InnoDB: Page dump in ascii and hex (16384 bytes):
```



数，造，未，来



# Partial Writes模拟：MySQL篇总结



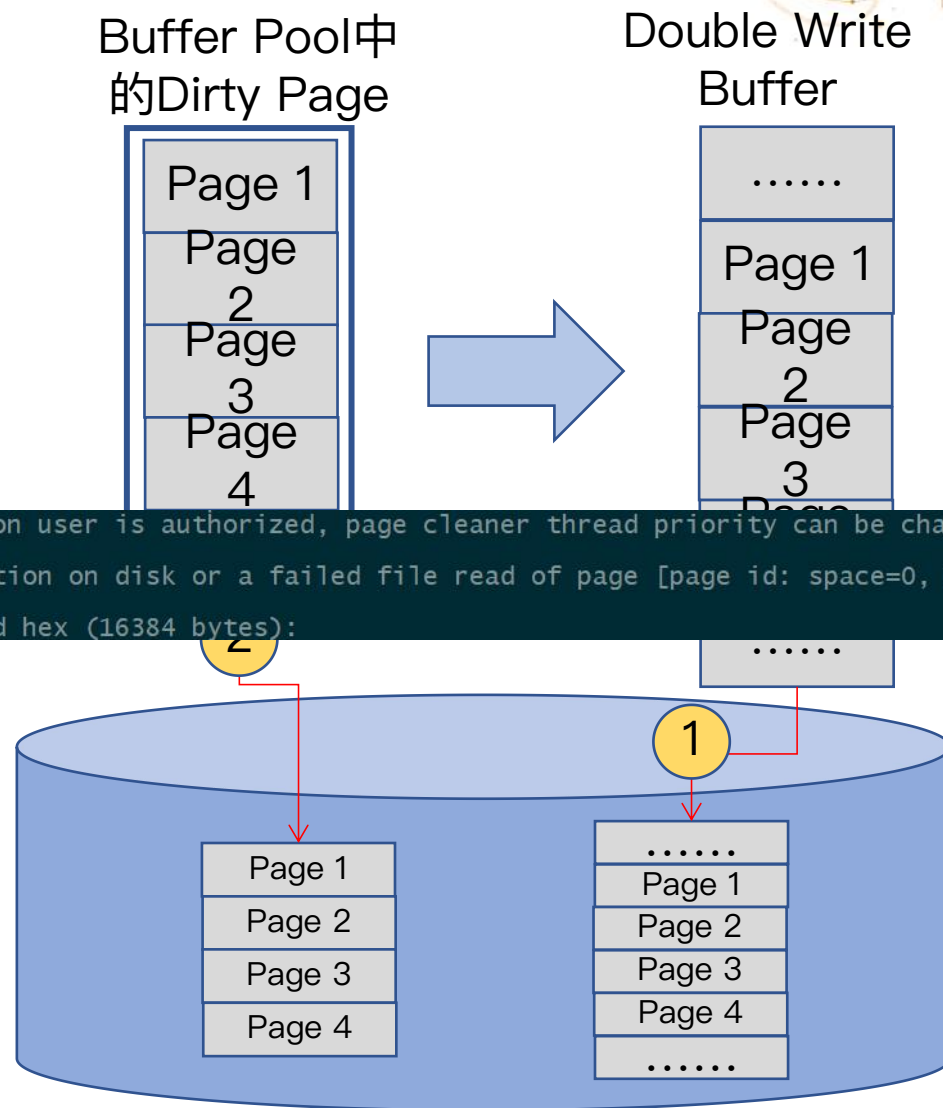
DTCC 2021  
第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021

Double Write并没有像想像的那样，发挥它应有的作用。

备份很重要。双写只在部分情况下有效。

解决Partial Writes，终极方案，如myerror.log中所说：

```
2021-10-09T15:20:31.779551+08:00 0 [Note] InnoDB: If the mysqld execution user is authorized, page cleaner thread priority can be changed. See the man page of setpriority().  
2021-10-09T15:20:31.797777+08:00 0 [ERROR] InnoDB: Database page corruption on disk or a failed file read of page [page id: space=0, page number=5]. You may have to recover from a backup.  
2021-10-09T15:20:31.797831+08:00 0 [Note] InnoDB: Page dump in ascii and hex (16384 bytes):
```



数，造，未，来



# Partial Writes模拟：Oracle篇

➤ I/O相关系统函数（和MySQL一样）

打开异步I/O：Sys\_io\_submit

未打开异步I/O：pwrite







## ➤ 准备测试数据

```
create tablespace TPS_TEST datafile '/u01/oradata/PROD/tps_test_01' size 128k;
```

```
create table u2.dtcc(id1 int primary key, c1 varchar2(30)) tablespace tps_test;
```

```
insert into u2.dtcc values(1, 'AAAAAAAAA');
```

```
commit;
```

Oracle 是从下往上使用页空间，虽然只插入一行，但目标行在页的尾部。




# Partial Writes模拟：Oracle篇



DTCC 2021  
第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021

## ➤ 准备测试数据



## ➤ 测试步骤

- 修改目标行，使目标块成为脏页。Oracle会修改块头部的SCN等信息，和块尾部的目标行。
- 手动发出检查点，触发DBWR进程写脏块。
- 拦截I/O系统调用，修改count参数，造成一个8K块，前4K写成功、后4K写失败。
- 观察、记录Oracle的处理过程。



数，造，未，来



IT168.com

ChinaUnix.net

ITPUB

# Partial Writes模拟：Oracle篇



**DTCC 2021**  
第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021

## ➤ 首先观察I/O:

```
[root@localhost iotest]# stap -g dbpage_o.stp ora_dbw0_prod
Begin...
20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 102 28aeb6000 2000 19b2000
return:20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 103 29fc8c000 2000 17b40000
return:20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 103 29271c000 2000 17b46000
return:20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 104 28d8d4000 2000 21878000
return:20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 104 28cf8c000 2000 251a8000
return:20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 104 28b26c000 2000 25202000
return:20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 104 28cf6a000 2000 25284000
return:20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 104 28c898000 2000 252e2000
return:20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 104 28c84e000 2000 25478000
return:20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
```



数，造，未，来







## ➤ 首先观察I/O:

```
[root@localhost iotest]# stap -g dbpage_o.stp ora_dbw0_prod
Begin...
20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 102 28aeb6000 2000 19b2000
return:20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 103 29fc8c000 2000 17b40000
return:20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 103 29271c000 2000 17b46000
return:20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 104 28d8d4000 2000 21878000
return:20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 104 28cf8c000 2000 251a8000
return:20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 104 28b26c000 2000 25202000
return:20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 104 28cf6a000 2000 25284000
return:20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 104 28c898000 2000 252e2000
return:20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 104 28c84e000 2000 25478000
return:20177 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
```

0x102是system01.dbf, 0x103是undotbs01.dbf, 0x104是sysaux01.dbf。后面还有0x105, undotbs01\_02.dbf文件。0x108, 是目标文件tps\_test\_01。







## ➤ 开始测试

```
probe syscall.pwrite
{
    if (execname() == @1)
    {
        printf("%d %s %s %x %x %x %x\n", tid(), pp(), ppfunc(), $fd, $buf, $count, $pos)
        if ( ($fd == 0x102 || $fd == 0x103 || $fd == 0x104 || $fd == 0x105 || $fd == 0x108) && ($count = 0x2000) )
        {
            $count = 0x1000
            printf("Update Value\n")
        }
        target = 1
    }
}
```

对部分文件的\$count参数进行修改。



# Partial Writes模拟：Oracle篇

## ➤ 开始测试

```
[root@localhost ~]# sta
Begin...
20177 kernel.function("Sys_p
Update Value
return:20177 kernel.function
20177 kernel.function("Sys_p
Update Value
return:20177 kernel.function
20177 kernel.function("Sys_p
Update Value
return:20177 kernel.function
20177 kernel.function("Sys_p
Update Value
return:20177 kernel.function
20177 kernel.function("Sys_p
```

```
SQL> update u2.dtcc set c1=lower(c1) where id1=1;
commit;
alter system checkpoint;
```

已更新 1 行。

0 2000 20f5c000

```
SQL>
```

提交完成。

0 2000 25190000

```
SQL>
```

```
alter system checkpoint
```

0 2000 251a8000

第 1 行出现错误:

0 2000 25236000

ORA-03113: 通信通道的文件结尾 进程 ID:

22538

0 2000 25478000

会话 ID: 21262 序列号: 20043





## ➤ 开始测试

```
SQL> startup mount;  
ORACLE 例程已经启动。  
  
Total System Global Area 9663673928 bytes  
Fixed Size                  8906312 bytes  
Variable Size              2550136832 bytes  
Database Buffers           7079985152 bytes  
Redo Buffers                24645632 bytes  
  
数据库装载完毕。  
SQL> SQL> SQL> alter database open;  
alter database open  
*  
第 1 行出现错误:  
ORA-00603: ORACLE server session terminated by fatal error  
ORA-01092: ORACLE instance terminated. Disconnection forced  
ORA-01578: ORACLE data block corrupted (file # 4, block # 3664)  
ORA-01110: data file 4: '/u01/oradata/PROD/undotbs01_02.dbf'  
进程 ID: 26127  
会话 ID: 6242 序列号: 12926
```







## ➤ 开始测试

发现非正常关库，开始实例恢复

```
2021-10-09T17:10:42.286944+08:00
alter database open
2021-10-09T17:10:42.345428+08:00
Ping without log force is disabled:
  instance mounted in exclusive mode.
Buffer Cache Full DB Caching mode changing from FULL CACHING DISABLED to FULL CACHING ENABLED
2021-10-09T17:10:42.416860+08:00
Beginning crash recovery of 1 threads
  parallel recovery started with 3 processes
  Thread 1: Recovery starting at checkpoint rba (logseq 7112 block 1736), scn 0
2021-10-09T17:10:42.797432+08:00
Started redo scan
2021-10-09T17:10:42.850287+08:00
Completed redo scan
  read 74 KB redo, 24 data blocks need recovery
2021-10-09T17:10:42.920348+08:00
Hex dump of (file 3, block 75976) in trace file /u01/diag/rdbms/prod/prod/trace/prod_p000_26219.trc

Corrupt block relative dba: 0x00c128c8 (file 3, block 75976)
Fractured block found during crash/instance recovery
Data in bad block:
  type: 32 format: 2 rdba: 0x00c128c8
  last change scn: 0x0000.0000.06c3c72d seq: 0x1 flg: 0x04
  spare3: 0x0
  consistency value in tail: 0xbdbd2001
  check value in block header: 0xa7f3
  computed block checksum: 0x7990
```







## ➤ 开始测试

确定恢复的起点：

- 7112号Redo文件
- 1736号块处

```
2021-10-09T17:10:42.286944+08:00
alter database open
2021-10-09T17:10:42.345428+08:00
Ping without log force is disabled:
  instance mounted in exclusive mode.
Buffer Cache Full DB Caching mode changing from FULL CACHING DISABLED to FULL CACHING ENABLED
2021-10-09T17:10:42.416860+08:00
Beginning crash recovery of 1 threads
  parallel recovery started with 3 processes
  Thread 1: Recovery starting at checkpoint rba (logseq 7112 block 1736), scn 0
2021-10-09T17:10:42.797432+08:00
Started redo scan
2021-10-09T17:10:42.850287+08:00
Completed redo scan
  read 74 KB redo, 24 data blocks need recovery
2021-10-09T17:10:42.920348+08:00
Hex dump of (file 3, block 75976) in trace file /u01/diag/rdbms/prod/prod/trace/prod_p000_26219.trc

Corrupt block relative dba: 0x00c128c8 (file 3, block 75976)
Fractured block found during crash/instance recovery
Data in bad block:
  type: 32 format: 2 rdba: 0x00c128c8
  last change scn: 0x0000.0000.06c3c72d seq: 0x1 flg: 0x04
  spare3: 0x0
  consistency value in tail: 0xbabd2001
  check value in block header: 0xa7f3
  computed block checksum: 0x7990
```




# Partial Writes模拟：Oracle篇



DTCC 2021

第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021



Record 1 Record 2 .....



数，造，未，来

Hadoop

OLAP

HANA

MPP

DB2

Oracle

SQL




# Partial Writes模拟：Oracle篇



DTCC 2021

第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021



数，造，未，来

Hadoop

OLAP

HANA

MPP

DB2

Oracle

SQL

IT168.com

ChinaUnix.net


ITPUB



# Partial Writes模拟：Oracle篇



DTCC 2021  
第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021



Record 1 Record 2 .....



数，造，未，来

Hadoop

OLAP

HANA

MPP

DB2

Oracle

SQL

IT168.com

ChinaUnix.net

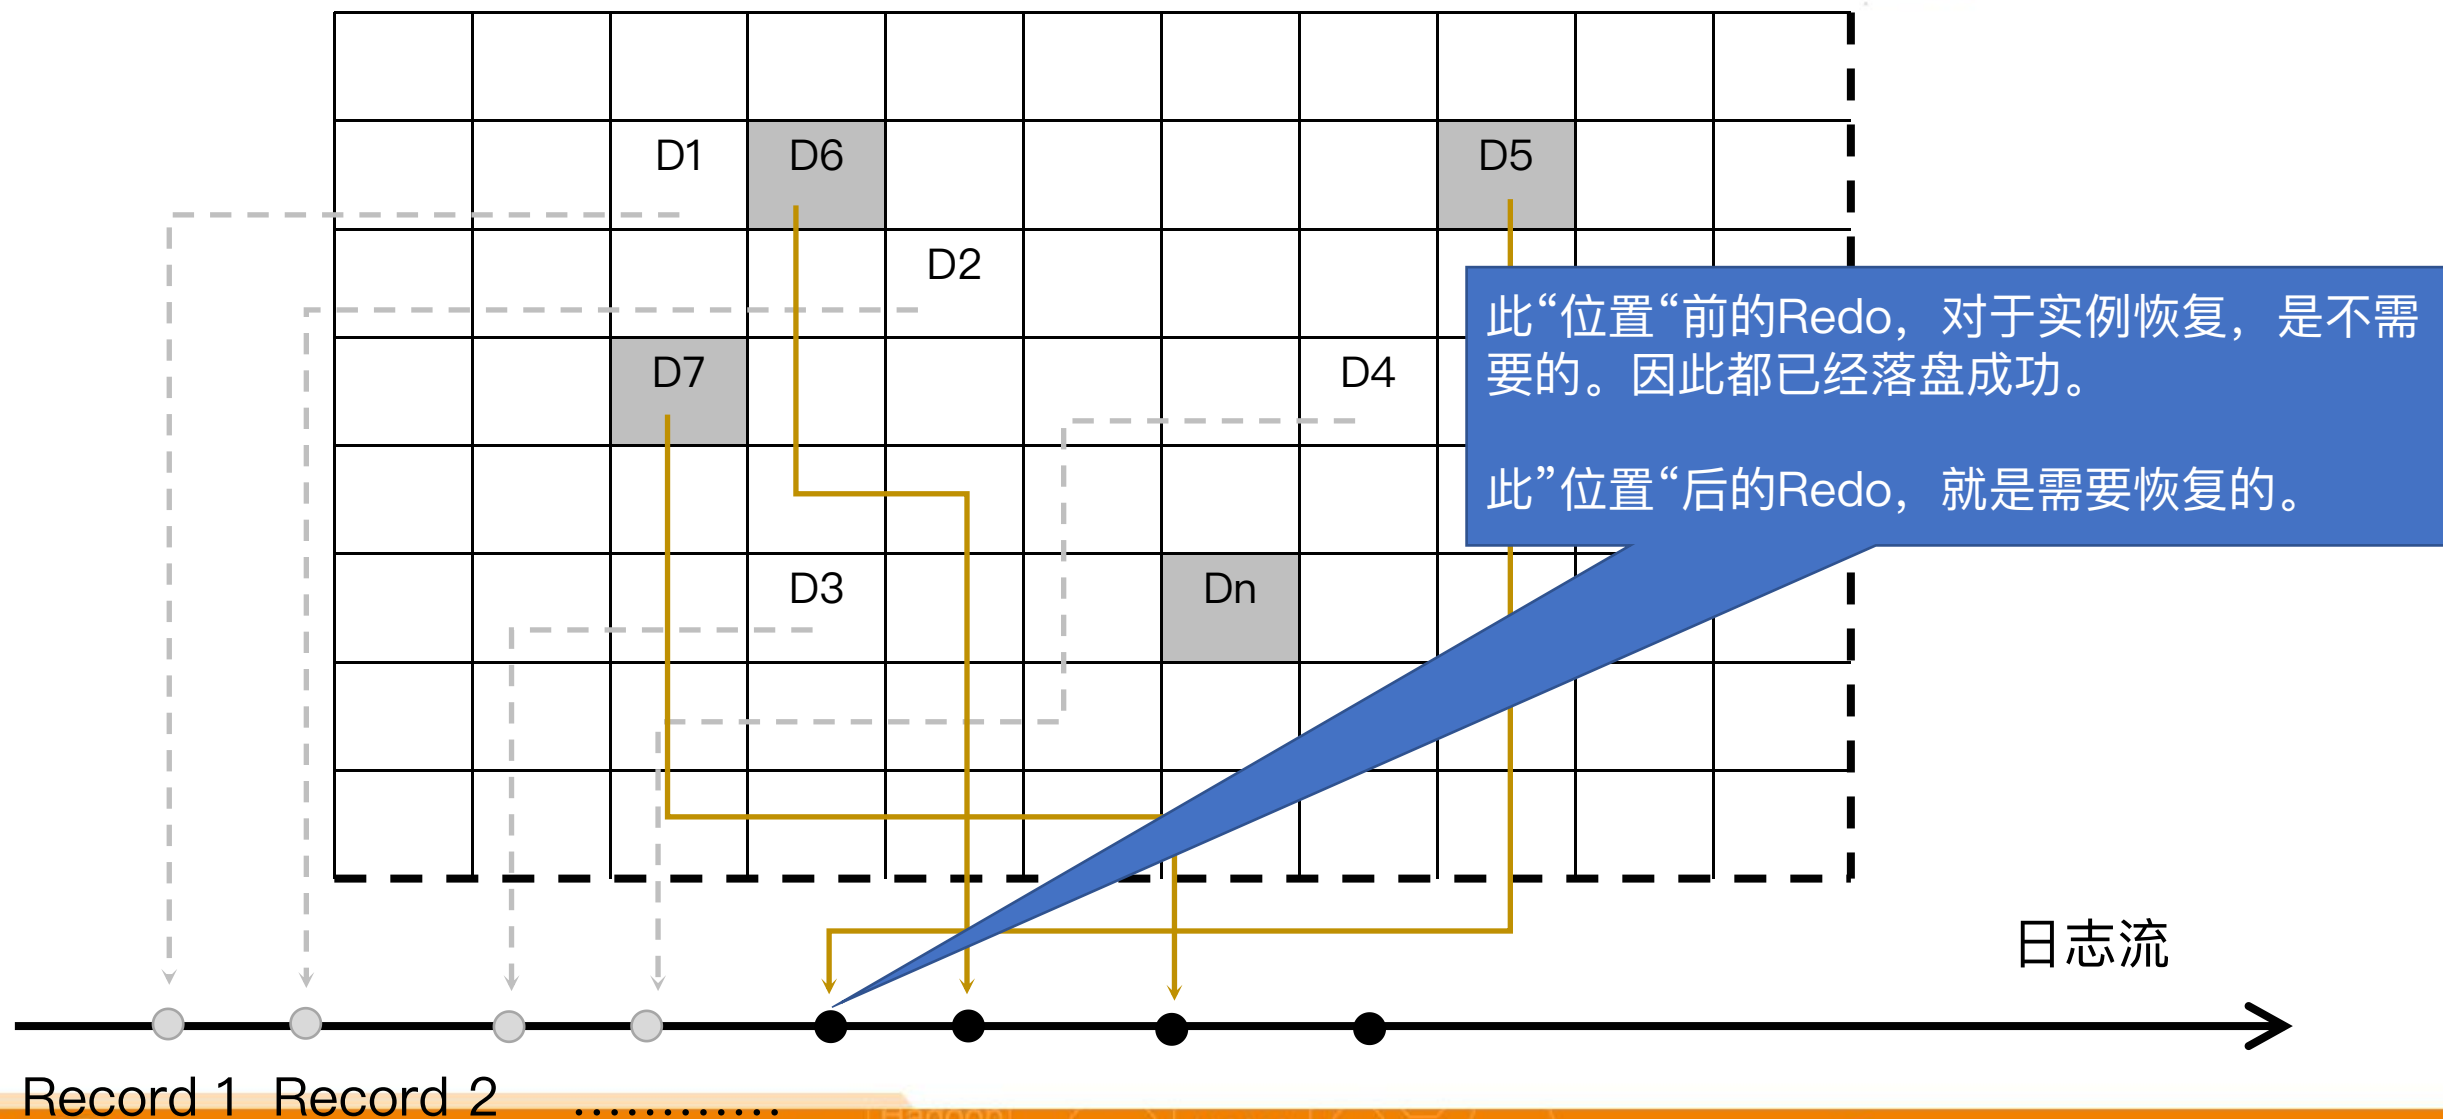
ITPUB

# Partial Writes模拟：Oracle篇



DTCC 2021

第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021



Record 1 Record 2 .....



数，造，未，来

Hadoop

OLAP

HANA

MPP

Oracle

SQL

IT168.com

ChinaUnix.net

ITPUB

# Partial Writes模拟：Oracle篇



DTCC 2021  
第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021

## ➤ 开始测试

```
2021-10-09T17:10:42.286944+08:00
alter database open
2021-10-09T17:10:42.345428+08:00
Ping without log force is disabled:
  instance mounted in exclusive mode.
Buffer Cache Full DB Caching mode changing from FULL CACHING D
2021-10-09T17:10:42.416860+08:00
Beginning crash recovery of 1 threads
  parallel recovery started with 3 processes
  Thread 1: Recovery starting at checkpoint rba (logseq 7112 block 1736), scn 0
2021-10-09T17:10:42.797432+08:00
Started redo scan
2021-10-09T17:10:42.850287+08:00
Completed redo scan
  read 74 KB redo, 24 data blocks need recovery
2021-10-09T17:10:42.920348+08:00
Hex dump of (file 3, block 75976) in trace file /u01/diag/rdbms/prod/prod/trace/prod_p000_26219.trc

Corrupt block relative dba: 0x00c128c8 (file 3, block 75976)
Fractured block found during crash/instance recovery
Data in bad block:
  type: 32 format: 2 rdba: 0x00c128c8
  last change scn: 0x0000.0000.06c3c72d seq: 0x1 flg: 0x04
  spare3: 0x0
  consistency value in tail: 0xbdbd2001
  check value in block header: 0xa7f3
  computed block checksum: 0x7990
```

确定恢复起点后，扫描Redo文件，判断出，有24个脏块需要恢复。



数，造，未，来



IT168.com

ChinaUnix.net

ITPUB



# Partial Writes模拟：Oracle篇



DTCC 2021  
第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021

## ➤ 开始测试

```
2021-10-09T17:10:42.286944+08:00
alter database open
2021-10-09T17:10:42.345428+08:00
Ping without log force is disabled:
  instance mounted in exclusive mode.
Buffer Cache Full DB Caching mode changing from FULL CACHING D
2021-10-09T17:10:42.416860+08:00
Beginning crash recovery of 1 threads
  parallel recovery started with 3 processes
  Thread 1: Recovery starting at checkpoint rba (logseq 7112 b1
2021-10-09T17:10:42.797432+08:00
Started redo scan
2021-10-09T17:10:42.850287+08:00
Completed redo scan
  read 74 KB redo, 24 data blocks need recovery
2021-10-09T17:10:42.920348+08:00
Hex dump of (file 3, block 75976) in trace file /u01/diag/rdbms/prod/prod/trace/prod_p000_26219.trc

Corrupt block relative dba: 0x00c128c8 (file 3, block 75976)
Fractured block found during crash/instance recovery
Data in bad block:
  type: 32 format: 2 rdba: 0x00c128c8
  last change scn: 0x0000.0000.06c3c72d seq: 0x1 flg: 0x04
  spare3: 0x0
  consistency value in tail: 0xbdbd2001
  check value in block header: 0xa7f3
  computed block checksum: 0x7990
```

Partial Writes的坏块，无法恢复。



数，造，未，来

Hadoop

OLAP

HANA

MPP

DB2

Oracle

SQL

IT168.com

ChinaUnix.net

ITPUB



## ➤ 恢复的局限性

页（块）的损坏是千奇百怪的。数据又是十分重要的。

尝试用Redo在坏块基础上进行修复。坏块坏的可能近乎无限，你能想到的坏的情况是有限。





## ➤ 恢复的局限性

页（块）的损坏是千奇百怪的。数据又是十分重要的。

尝试用Redo在坏块基础上进行修复。坏块坏的可能近乎无限，你能想到的坏的情况是有限：

## 以有涯随无涯，殆已

因此，数据库的方式都是，在前一份完好一致的、正确的数据基础上，应用日志（redo、xlog、binlog等），将数据推进到最近的时刻。





# Partial Writes模拟：Oracle篇总结



**DTCC 2021**  
第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021

## ➤ Oracle篇总结

Oracle应对Partial Writes，依赖“检查”和介质恢复。



数，造，未，来




IT168.com

ChinaUnix.net

ITPUB

# Partial Writes模拟：Oracle篇总结





➤ I/O相关函数：pwrite

➤ 准备测试数据：

```
create table dtcc2(id1 int primary key, id2 int, c1 varchar(30), c2 varchar(30)) ;
```

```
insert into dtcc2 values(1, 101, 'aaaaa1', 'aaaaaa1');
```

和Oracle一样，PG也是从下往上使用页空间，虽然只插入一行，但目标行在页的尾部。

如下方式观察表对应的文件：

```
cd PG的数据库目录/base/DATABASE_OID (Database_OID来自于pg_database视图)  
hexdump -C TABLE_OID (Table_OID来自于pg_class视图)
```






# Partial Writes模拟：PG篇



DTCC 2021  
第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021

## ➤ 准备测试数据



## ➤ 测试步骤：和Oracle方法类似

- 修改目标行，使目标块成为脏页。
- 手动发出检查点，触发CheckPoint进程写脏块。
- 拦截I/O系统调用，修改count参数，造成一个8K块，前4K写成功、后4K写失败。
- 观察、记录PG的处理过程。



数，造，未，来



IT168.com

ChinaUnix.net

ITPUB

## ➤ 观察I/O:

```
probe syscall.pwrite
{
    if (execname() == @1)
    {
        printf("old:%d %s %s %s %x %x %x %x\n", tid(), execname(), pp(), ppfunc(), $fd, $buf,
$count, $pos)
    }
}
```

## ➤ 开始测试

```
postgres=# update dtcc2 set c1=upper(c1) where id1=1;
UPDATE 1
postgres=# checkpoint;
```



# Partial Writes模拟：PG篇



DTCC 2021  
第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021

➤ 观察I/O:

用户后台进程写WAL日志

probe syscall.pwrite

```
[root@localhost iotest]# stap -g dbpage_pg.stp postgres
Begin
old:7703 postgres kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 8 2aaaaac30000 2000 1a000
return:7703 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64

old:16047 postgres kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 3 2aaaaac30000 2000 1a000
return:16047 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64

old:16045 postgres kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 4 2aaaaf0d2980 2000 0
return:16045 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
old:16045 postgres kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 3 2aaaaac30000 2000 1a000
return:16045 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64

old:16047 postgres kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 3 2aaaaac30000 2000 1a000
return:16047 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
```

postgres=# checkpoint;



数，造，未，来





# Partial Writes模拟：PG篇



DTCC 2021  
第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021

➤ 观察I/O:

walwriter进程写WAL日志

probe syscall.pwrite

```
[root@localhost iotest]# stap -g dbpage_pg.stp postgres
Begin...
old:7703 postgres kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 8 2aaaaac30000 2000 1a000
return:7703 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64

old:16047 postgres kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 3 2aaaaac30000 2000 1a000
return:16047 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64

old:16045 postgres kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 4 2aaaaaf0d2980 2000 0
return:16045 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
old:16045 postgres kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 3 2aaaaac30000 2000 1a000
return:16045 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64

old:16047 postgres kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 3 2aaaaac30000 2000 1a000
return:16047 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
```

postgres=# checkpoint;



数，造，未，来



# Partial Writes模拟：PG篇

➤ 观察I/O:

Checkpoint进程写目标表对应的文件

probe syscall.pwrite

```
[root@localhost iotest]# stap -g dbpage_pg.stp postgres
Begin...
old:7703 postgres kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 8 2aa... 00 2000 1a000
return:7703 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64

old:16047 postgres kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 3 2... ac30000 2000 1a000
return:16047 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64

old:16045 postgres kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 4 2aaaaf0d2980 2000 0
return:16045 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
old:16045 postgres kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 3 2aaaaac30000 2000 1a000
return:16045 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64

old:16047 postgres kernel.function("Sys_pwrite64@fs/read_write.c:621").call Sys_pwrite64 3 2aaaaac30000 2000 1a000
return:16047 kernel.function("Sys_pwrite64@fs/read_write.c:621").return Sys_pwrite64
```

postgres=# checkpoint;



数，造，未，来





# Partial Writes模拟：PG篇



DTCC 2021  
第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021

## ➤ 测试脚本：

```
probe syscall.pwrite
{
    if (execname() == @1)
    {
        printf("old:%d %s %s %s %x %x %x %x\n", tid(), execname(), pp(), ppfunc(), $fd, $buf, $count, $pos)
        if ( (tid() == 16045 && $fd == 0x4) || bufaddr == $buf )
        {
            $count = 0x1000
            if (bufaddr == $buf)
            {
                $buf = $buf - (i * 0x1000)
                $pos = $pos - (i * 0x1000)
                $fd = 0x1234
                i = 0
            }
            bufaddr = $buf
            i += 1
        }
        printf("new:%d %s %s %s %x %x %x %x\n", tid(), execname(), pp(), ppfunc(), $fd, $buf, $count, $pos)
        target = 1
    }
}
```

uf,



数，造，未，来



IT168.com

ChinaUnix.net

ITPUB



# Partial Writes模拟：PG篇



DTCC 2021  
第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021

## ➤ 测试结果：

```
postgres=# update dtcc2 set c1=lower(c1) where id1=1;
UPDATE 1
postgres=# checkpoint;
ERROR:  checkpoint request failed
HINT:  Consult recent messages in the server log for details.
postgres=#
postgres=#
```

I/O错误后数据库没有当掉，使用kill命令，kill掉所有PG的进程，模拟意外当库：

```
postgres 5385 1 0 11:36 ? 00:00:00 /home/postgres/postgresql-12.1/prebuild/bin/postgres -D /data/pgdata
postgres 5386 5385 0 11:36 ? 00:00:00 postgres: logger
postgres 5388 5385 0 11:36 ? 00:00:00 postgres: checkpointer
postgres 5389 5385 0 11:36 ? 00:00:00 postgres: background writer
postgres 5390 5385 0 11:36 ? 00:00:00 postgres: walwriter
postgres 5391 5385 0 11:36 ? 00:00:00 postgres: autovacuum launcher
postgres 5392 5385 0 11:36 ? 00:00:00 postgres: archiver
postgres 5393 5385 0 11:36 ? 00:00:00 postgres: stats collector
postgres 5394 5385 0 11:36 ? 00:00:00 postgres: logical replication launcher
postgres 5399 2395 0 11:36 pts/2 00:00:00 psql -p6016
postgres 5400 5385 0 11:36 ? 00:00:00 postgres: postgres postgres [local] idle
root 5540 2547 99 11:38 pts/3 00:00:03 stap -g dbpage_pg.stp postgres
postgres 5544 2285 0 11:38 pts/0 00:00:00 ps -ef
postgres 5545 2285 0 11:38 pts/0 00:00:00 grep --color=auto postgres
```

kill -9 5385 5386 5388 5389 5390 5391 5392 5393 5394 5400



数，造，未，来



# Partial Writes模拟：PG篇

## ➤ 测试结果：

- 查看日志，有大量I/O错误：

```
2021-10-10 22:56:46.732 EDT,,2668,,6163a7ee.a6c,2,,2021-10-10 22:56:46 EDT,,0,LOG,00000,"database system is ready to accept connections",,,,,,""  
2021-10-10 23:34:57.654 EDT,,2671,,6163a7ee.a6f,1,,2021-10-10 22:56:46 EDT,,0,ERROR,53100,"could not write block 0 in file ""base/13593/34110"": wrote onl  
y 4096 of 8192 bytes",,"check free disk space.",,"writing block 0 of relation base/13593/34110",,,,,,""  
2021-10-10 23:34:57.654 EDT,"postgres","postgres",2680,["local"],6163a7f0.a78,1,"CHECKPOINT",2021-10-10 22:56:48 EDT,3/7,0,ERROR,XX000,"checkpoint request  
failed",,"Consult recent messages in the server log for details.",,"checkpoint;",,"psql"  
2021-10-10 23:34:59.706 EDT,,2671,,6163a7ee.a6f,2,,2021-10-10 22:56:46 EDT,,0,ERROR,XX000,"could not write block 0 in file ""base/13593/34110"": Invalid a  
rgument",,"writing block 0 of relation base/13593/34110",,,,,,""  
2021-10-10 23:34:59.706 EDT,,2671,,6163a7ee.a6f,3,,2021-10-10 22:56:46 EDT,,0,WARNING,58030,"could not write block 0 of base/13593/34110",,"Multiple failur  
es --- write error might be permanent.",,,,,,""  
2021-10-10 23:35:00.771 EDT,,2671,,6163a7ee.a6f,4,,2021-10-10 22:56:46 EDT,,0,ERROR,53100,"could not write block 0 in file ""base/13593/34110"": wrote onl  
y 4096 of 8192 bytes",,"check free disk space.",,"writing block 0 of relation base/13593/34110",,,,,,""  
2021-10-10 23:35:00.771 EDT,,2671,,6163a7ee.a6f,5,,2021-10-10 22:56:46 EDT,,0,WARNING,58030,"could not write block 0 of base/13593/34110",,"Multiple failur  
es --- write error might be permanent.",,,,,,""  
2021-10-10 23:35:01.820 EDT,,2671,,6163a7ee.a6f,6,,2021-10-10 22:56:46 EDT,,0,ERROR,XX000,"could not write block 0 in file ""base/13593/34110"": Invalid a  
rgument",,"writing block 0 of relation base/13593/34110",,,,,,""  
2021-10-10 23:35:01.820 EDT,,2671,,6163a7ee.a6f,7,,2021-10-10 22:56:46 EDT,,0,WARNING,58030,"could not write block 0 of base/13593/34110",,"Multiple failur  
es --- write error might be permanent.",,,,,,""
```

- 重新启动PG，PG自动进行恢复：

```
,2021-10-10 23:35:42 EDT,,0,LOG,00000,"database system was interrupted;",,,,,,""  
,2021-10-10 23:35:42 EDT,,0,LOG,00000,"database system was not properly shut down; automatic recovery",,,,,,""  
,2021-10-10 23:35:42 EDT,,0,LOG,00000,"redo starts at 1F/2E000310",,,,,,""  
,2021-10-10 23:35:42 EDT,,0,LOG,00000,"invalid record length at 1F/2E000688: wanted 24, got 0",,,,,,""  
,2021-10-10 23:35:42 EDT,,0,LOG,00000,"redo done at 1F/2E000650",,,,,,""  
,2021-10-10 23:35:42 EDT,,0,LOG,00000,"database system is ready to accept connections",,,,,,""  
,2021-10-10 23:35:42 EDT,,0,LOG,00000,"received fast shutdown request",,,,,,""  
,2021-10-10 23:35:42 EDT,,0,LOG,00000,"aborting any active transactions",,,,,,""
```

1F/2E000310，就是“检查点位置”  
相关信息记录在PG的控制文件中







## ➤ 测试结果：

```
[postgres@localhost ~]$ pg_ctl -D /data/pgdata -l logfile start
waiting for server to start.... done
server started
[postgres@localhost ~]$
[postgres@localhost ~]$
[postgres@localhost ~]$ psql -p6016
psql (12.1)
Type "help" for help.

postgres=# select * from dtcc2;
 id1 | id2 |  c1  |  c2
-----+-----+-----+-----
   1 | 101 | aaaaa1 | aaaaaa1
(1 row)

postgres=#
```

再次启动数据库，一切正常，数据也没有丢失。最后的测试语句，就是将数据改为小写。







## ➤ 测试结果：

```
[postgres@localhost ~]$ pg_ctl -D /data/pgdata -l logfile start  
waiting for server to start.... done  
server started  
[postgres@localhost ~]$
```

- 对数据库而言，Partial Writes是偶而出现的情况，操作系统不会频繁的崩溃。
- 为了一个极其偶然出现的Partial Writes，引入极大性能损耗的FPW，值得商榷。毕竟，像Oracle这样成熟的商业数据库，并没有类似FPW的方式。

```
postgres=#
```

再次启动数据库，一切正常，数据也没有丢失。最后的测试语句，就是将数据改为小写。

PG以巨大的性能代价（超30%性能下降），解决了Partial Writes问题。





## ➤ 测试结果：

测试步骤相同，唯一区别FPW参数设置为OFF：

- 执行脚本
- 发出Update SQL和Checkpoint命令
- kill 掉所有PG进程
- 重启数据库，观察结果



# Partial Writes模拟：PG篇 -- 关闭FPW的测试



DTCC 2021

大会  
2021

```
postgres=# update dtcc2 set c1=lower(c1) where id1=1;  
checkpoint;
```

```
UPDATE 1
```

```
postgres=# checkpoint;  
ERROR: checkpoint request failed  
HINT: Consult recent messages in the server log for details.  
postgres=#  
postgres=# select * from dtcc2;  
server closed the connection unexpectedly  
This probably means the server terminated abnormally  
before or while processing the request.  
The connection to the server was lost. Attempting reset: Failed.
```

```
!> \q  
[postgres@localhost ~]$  
[postgres@localhost ~]$ psql -p6016  
psql: error: could not connect to server: could not connect to server: Connection refused  
Is the server running locally and accepting  
connections on Unix domain socket "/tmp/.s.PGSQL.6016"?
```

```
[postgres@localhost ~]$  
[postgres@localhost ~]$  
[postgres@localhost ~]$ pg_ctl -D /data/pgdata -l logfile start  
pg_ctl: another server might be running; trying to start server anyway  
waiting for server to start..... done  
server started  
[postgres@localhost ~]$ psql -p6016  
psql (12.1)  
Type "help" for help.
```

```
postgres=# select * from dtcc2;  
 id1 | id2 |  c1  |  c2  
-----+-----+-----+-----  
   1 | 101 | AAAA1 | aaaaa1  
(1 row)
```

```
postgres=#
```

Update已经完成（隐含提交），  
事务已经完成，C1列将被改为小写。



# Partial Writes模拟：PG篇 -- 关闭FPW的测试

```
postgres=# update dtcc2 set c1=lower(c1) where id1=1;  
checkpoint;
```

```
UPDATE 1
```

```
postgres=# checkpoint;  
ERROR:  checkpoint request failed  
HINT:  Consult recent messages in the server log for details.  
postgres=#
```

```
postgres=# select * from dtcc2;  
server closed the connection unexpectedly  
This probably means the server terminated abnormally  
before or while processing the request.  
The connection to the server was lost. Attempting reset: Failed.
```

```
!> \q  
[postgres@localhost ~]$  
[postgres@localhost ~]$ psql -p6016  
psql: error: could not connect to server: could not connect to server: Connection refused  
Is the server running locally and accepting  
connections on Unix domain socket "/tmp/.s.PGSQL.6016"?
```

```
[postgres@localhost ~]$  
[postgres@localhost ~]$  
[postgres@localhost ~]$ pg_ctl -D /data/pgdata -l logfile start  
pg_ctl: another server might be running; trying to start server anyway  
waiting for server to start..... done  
server started  
[postgres@localhost ~]$ psql -p6016  
psql (12.1)  
Type "help" for help.
```

```
postgres=# select * from dtcc2;  
 id1 | id2 |  c1  |  c2  
-----+-----+-----+-----  
   1 | 101 | AAAA1 | aaaaaa1  
(1 row)
```

```
postgres=#
```

Checkpoint失败，写数据错误，出现Partial Writes。

# Partial Writes模拟：PG篇 -- 关闭FPW的测试



DTCC 2021

第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021

```
postgres=# update dtcc2 set c1=lower(c1) where id1=1;
checkpoint;

UPDATE 1
postgres=# checkpoint;
ERROR:  checkpoint request failed
HINT:  Consult recent messages in the server log for details.
postgres=#
postgres=# select * from dtcc2;
server closed the connection unexpectedly
        This probably means the server terminated abnormally
        before or while processing the request.
The connection to the server was lost. Attempting reset: Failed.
!> \q
[postgres@localhost ~]$
[postgres@localhost ~]$ psql -p6016
psql: error: could not connect to server: could not connect to server: Connection refused
        Is the server running locally and accepting
        connections on Unix domain socket "/tmp/.s.PGSQL.6016"?
[postgres@localhost ~]$
[postgres@localhost ~]$
[postgres@localhost ~]$ pg_ctl -D /data/pgdata -l logfile start
pg_ctl: another server might be running; trying to start server anyway
waiting for server to start..... done
server started
[postgres@localhost ~]$ psql -p6016
psql (12.1)
Type "help" for help.

postgres=# select * from dtcc2;
 id1 | id2 |  c1  |  c2
-----+-----+-----+-----
   1 | 101 | AAAA1 | aaaaaa1
(1 row)
```

重启数据库后，最后已经提交的更新，丢失。数据仍是大写'A'。

# PostgreSQL Partial Writes解决方案的优化



DTCC 2021  
第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021

## ➤ 为什么不能参照MySQL的Double Write（双写）

- 前文已经有测试，无法恢复所有Partial Writes场景
- 根据前文关掉FPW的测试，数据存在不一致的风险

## ➤ 参照

**结论：PG中关闭FPW引入双写，存在数据不一致的可能**

- 报出错误
- 最终依赖基于备份的恢复



数，造，未，来



IT168.com

ChinaUnix.net

ITPUB



# PostgreSQL Partial Writes解决方案的优化



DTCC 2021  
第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021

## ➤ 摸着Oracle过河

- PG和Oracle使用一样的物理日志，备份恢复体系的原理一致
- Oracle在控制文件中记录每次检查点完成后的检查点位置，PG也有一模一样的机制。

## ➤ 目标

- PG完成可以实现和Oracle一模一样的机制：在XLOG中，检查“检查点位置”后XLOG Record对应的块，如没被写坏，使用对应Record进行恢复。如已经被Partial Writes，写成坏块，报错。
- 安全的关闭FPW，（至少）提升30%性能。



数，造，未，来



IT168.com

ChinaUnix.net

ITPUB

# Partial Writes解决方案的优化



## ➤ PG的解决方案：

- PG的恢复在StartupXLOG()函数中。  
(xlog.c文件)
- 在此if之后，就是PG的恢复流程。

```
/* REDO */
if (InRecovery)
{
    int                rmid;

    /*
     * Update pg_control to show that we are recovering and to show the
     * selected checkpoint as the place we are starting from. We also mark
     * pg_control with any minimum recovery stop point obtained from a
     * backup history file.
     */
    dbstate_at_startup = ControlFile->state;
    if (InArchiveRecovery)
        ControlFile->state = DB_IN_ARCHIVE_RECOVERY;
    else
    {
        ereport(LOG,
                (errmsg("database system was not properly shut down; "
```



# Partial Writes解决方案的优化

## ➤ PG的解决方案：

- PG的恢复在StartupXLOG()函数中。  
(xlog.c文件)

- 在此if之后，就是PG的恢复流程。

- 目标：

并不直接修改StartupXLOG()

复制它，另外做一个Non-FPW模块，

如果出现异常当库，先使用Non-FPW模块进行检查、恢复，之后再使用正常的PG启动数据库。

```
/* REDO */
if (InRecovery)
{
    int                rmid;

    /*
     * Update pg_control to show that we are recovering and to show the
     * selected checkpoint as the place we are starting from. We also mark
     * pg_control with any minimum recovery stop point obtained from a
     * backup history file.
     */
    dbstate_at_startup = ControlFile->state;
    if (InArchiveRecovery)
        ControlFile->state = DB_IN_ARCHIVE_RECOVERY;
    else
    {
        ereport(LOG,
                (errmsg("database system was not properly shut down; "
```





# 欢迎讨论 共同为国内开源社区建设贡献力量



**DTCC 2021**  
第十二届中国数据库技术大会  
DATABASE TECHNOLOGY CONFERENCE CHINA 2021



**VAGE\_LV**



数 / 造 / 未 / 来





# THANKS